



MongoDB World  
2022


# *CSFLE and Queryable Encryption*



Pierre Petersson  
Advisory Solutions Architect EMEA

# Queryable Encryption and Client-Side FLE

## Design Goal



Provide some of the  **strongest levels of data privacy and security for regulated workloads**

- End to end protection of your most sensitive data
- Increase your confidence in moving to managed services in the cloud
- **Enable querying on encrypted data, in a secure, fast and scalable data (range, prefix)**
- Simplify compliance with modern privacy regs: “Right to Erasure

# Agenda

Encryption of Today

Encryption in Multi Cloud Environments

CSFLE vs Queryable Encryption

Key Management Enhancements

Q & A



Encryption  
is proven  
and trusted,  
but there are gaps




# Moving and storing data: most databases have it covered



*Reminder: TLS is to protect against network eavesdropping*




# Moving and storing data: most databases have it covered




*Reminder: At-rest encryption is (mostly) to protect non-running databases & backups*

# But what about data here?




# Data is in plain text while its being processed by the database



Data is vulnerable to insider access and active database breaches:

- Authorized and compromised administrators, DBAs & privileged users
- RAM scraping
- Process inspection
- Cloud providers

# Data is in plain text while its being processed by the database



This is why we built Client-Side Field Level Encryption!

Data is vulnerable to insider access and active database breaches:

- Authorized and compromised administrators , DBAs & users
- RAM scraping
- Process inspection



# How much of a risk is this...really?

**55%**

Of breaches committed by  
organized criminal groups,  
including state actors

**37%**


Of breaches used stolen  
credentials to access  
secure systems

**30%**

Of breaches  
were committed by  
internal users

**verizon**✓

[Verizon Data Breach  
Investigations Report, 2020](#)



- 
1. Do nothing and take the risk of unencrypted data in-use, in a live database **being exploited** and the company suffering a data breach
  2. Implementing **application level encryption** which adds **huge complexity and friction** to development processes

prag.

SUPERDUPER ENGLISH IDIOMS



# Application Encryption

Encrypt sensitive fields in the application, **before** it's sent to the database

## Code with no encryption

```
db.patients.insert({  
    firstName: 'Jane',  
    lastName: 'Doe',  
    ssn: "901-01-0002",  
    dob: new Date('1989-06-23'),  
})
```



## Code with explicit encryption

```
db.patients.insert({  
    firstName: 'Jane',  
    lastName: 'Doe',  
    ssn: db.getMongo().encrypt( key1 ,  
    "901-01-0002" , ENC_DETERM ),  
    dob: db.getMongo().encrypt( key1 ,  
    new Date('1989-06-23') , ENC_RANDOM)  
})
```



## CHALLENGES

# Application/Explicit Level Encryption



**Highly complex:** slows down application development



**Limits** application functionality, entire records returned as blobs, **not possible to search**



**Compromises** user experience requires additional client code

No database offers a standard solution to closing this gap...so what have we had to do?





# Automatic Encryption

## Code with no encryption

```
db.patients.insert({  
  firstName: 'Jane',  
  lastName: 'Doe',  
  ssn: "901-01-0002",  
  dob: new Date('1989-06-23'),  
})
```



## Code with automatic encryption

```
db.patients.insert({  
  firstName: 'Jane',  
  lastName: 'Doe',  
  ssn: "901-01-0002",  
  dob: new Date('1989-06-23'),  
})
```

No difference



# Automatic Encryption

## Code with no encryption

```
db.patients.insert({  
  firstName: 'Jane',  
  lastName: 'Doe',  
  ssn: "901-01-0002",  
  dob: new Date('1989-06-23'),  
})
```



## Code with explicit encryption

```
db.patients.insert({  
  firstName: 'Jane',  
  lastName: 'Doe',  
  ssn: "901-01-0002",  
  dob: new Date('1989-06-23'),  
})
```


## How?

Automatic encryption, available in MongoDB Atlas and Enterprise, uses [JSON Schema](#)


- Centrally defines which fields to encrypt and the algorithm to use



# Benefits of Automatic Encryption



- **Higher productivity:** No additional code for developers to write, reduces chance of human error
- **Improved data governance:** Centrally enforce the encryption of sensitive fields



# MongoDB Client Side FLE gives you a **much safer and simpler approach**

Select an encryption key, configure the fields to be encrypted in the MongoDB driver...and **GO**

- Sensitive data **never leaves the application** without first being encrypted. Data **remains encrypted** server-side
- No need to **modify applications**
- Encrypted data is **still queryable, equality match**
- **Minimal performance impact** to the database and the application
- **Reduce Cognitive load**, intuitive and easy for developers to configure and setup



# What about CSFLE in a **Multi Cloud** Environments



# Is Multi Cloud adopted by enterprises ?

60%

Small business  
(<100 Employees)

76%

Mid size companies  
(101-5000 Employees)


90%

Large enterprises



[HashiCorp State of Cloud Strategy Survey:](#)

# Challenges using CSFLE in a Multi Cloud Environment




# CSFLE in Multi Cloud Environments



# MongoDB Client Side FLE with KMIP



Allows Centralized, Cloud agnostic key management with KMIP




# Introducing Queryable Encryption


# Queryable Encryption



- Encrypt the sensitive data (fields)
  - Easy development cycle
  - No crypto experience required
  - Encrypted throughout the data lifecycle
  - **Rich expressive queries**
- MongoDB is the only platform to implement fast searchable encryption scheme
  - Server-side processing of encrypted data
  - Server does not know anything about the data



# Let's look closer



Encrypted fields are always stored, transmitted, processed, and retrieved as ciphertext, including queries

# Use Cases



## Industry: Financial Services

Bank application needs to find transactions using a **range of dates or dollar amounts** for fraud detection



## Industry: Human Resources

HR system allows searching for employees by the **last 4 digits** of their social security number



## Industry: Health Care

Customer support agents needs to find patient records by searching for the **first few characters** of their name



# Queryable Encryption – Key Benefits

## Reduce institutional risk

*Confident in storing and processing your sensitive workloads in MongoDB Atlas (Cloud)*

## Faster app development

*No crypto experience required*

*Intuitive and easy for developers to set up and use*



## Strong technical controls for critical data privacy use cases

*Meet the strictest data privacy requirements for confidentiality on security critical workloads*

## Ground-breaking query technology, standards-based cryptography

*Based on strong, standards-based cryptographic primitives*

## Rich querying on encrypted data

*Run expressive queries like range, equality, prefix, suffix, substring, and more on encrypted data*

## End-to-end fully randomized encryption

*Data never exists in the clear outside of the client  
Dramatically reduces attack surface*

# Our journey & roadmap



# Our Journey & Roadmap



Seny  
Kamara



Tarik  
Moataz

**Aroki Systems acquisition**  
Pioneers in Encrypted Search

2019

2021

May 2022

June 2022

Post 6.0

Post 6.0

Future

**Client-Side Field Level  
Encryption (CSFLE)**  
Equality search on  
Deterministic encryption

**Formation of Advanced  
Cryptography Research group**  
Seny Kamara, Tarik Moataz, and a  
team of PhD cryptography  
researchers

## Queryable Encryption Preview

Structured Encryption core  
functionality; *Equality search* on  
randomized encryption

## Queryable Encryption v1.2

Addition of *prefix*,*suffix*,  
*substring* query capabilities

**Queryable Encryption v1.1**  
Addition of *Range* query  
capabilities

New privacy-enhancing  
cryptography capabilities



# What is in 6.0 Public Preview?



## Crypto framework

Foundational work that enables equality and future query types



## Equality

Equality comparisons on randomly encrypted data



## Compass

Queryable Encryption Configuration Decryption



# What is a Public Preview?

- Available with 6.0 RC release
  - Evaluation only
  - May be breaking changes
  - Not recommended for production workloads




# Post 6.0 - Additional Query Types



Range



Prefix/Suffix



Substring

# CSFLE vs Queryable Encryption





# Both Are Supported

**CSFLE is NOT being deprecated**

**CSFLE to Queryable Encryption migration  
not (yet) supported**

CSFLE workloads should stay on CSFLE

**Queryable Encryption net new only**

Must be specified at collection creation



# CSFLE

- Client-side encryption
- Server is (largely) unaware
- Queryability
  - Equality only - Deterministic
    - Data leakage on low entropy fields
- Flexible key usage
  - unique key per field
  - 1 key for all fields
  - per-document keys
- No additional data elements

# Queryable Encryption

- Client-side encryption
- Server is integral
- Queryability
  - New functional search index
  - Equality - Fully random
    - No snapshot leakage, even on low entropy fields
  - Range, prefix, suffix and substring
- Requires a unique key per field
- Additional data
  - 1 new field per document
    - `__safeContent__`
  - 3 new system collections: `enxcol_.*`
  - **Do not modify any of these!**



# CSFLE vs Queryable Encryption Trade-offs

	Inserts	Find (equality)	Find (range, prefix, suffix, substring)	Storage Overhead	Frequency Leakage
FLE	Fast	Fast	No	Minimal	Possibly
Queryable Encryption	Slower	Fast	Yes	Yes	None

# Key Management Enhancements

(Coming Soon)






# Key Rotation




- Rotate the entire MongoDB key vault
- Previously, only new data encryption keys were protected by a new CMK
- Key Vault rotation replaces all former versions of CMK seamlessly, via a single API call



# Key Migration



- Changing from one key provider to another used to require decrypting and re-encrypting all of your data
- A single API call now seamlessly migrates your keys from any supported key provider to another one
  - AWS - GCP
  - Local - Azure
  - GCP - KMIP
- With no impact to your application or data



# Additional Resources

## **Resources CSFLE**

[Client-Side Field Level Encryption The Next Generation of Privacy & Security MDBW22 Video](#)

[Whitepaper CSFLE](#)

[CSFLE Multi Cloud Environments MDBW22 Video](#)

## **Queryable Encryption Resources**

Docs (very much a WIP): <https://www.mongodb.com/docs/upcoming/core/queryable-encryption/>

Blog post: <https://www.mongodb.com/blog/post/mongodb-releases-queryable-encryption-preview>

Product page & FAQ: <https://mongodb.com/products/queryable-encryption>



Thank you!

# When to use **Queryable** Encryption

When you need to do rich queries and not just equality match

If requirements are:

**Range queries**

E.g Date range

**Substring Queries**

E.g query part of string

**Prefix Queries**

E.g query on starts with

**Suffix Queries**

E.g query on ends with

**No frequency leakage**

Randomized ciphertext more secure





# When to use CSFLE Encryption

If requirements are:

## Fast inserts

Only one collection needs to be updated in contrast to 3 with Queryable Encryption

## Flexible key usage

One Key per field, one key per collection or combination.

## Only Equality match

E.g query on exact match using deterministic algorithm

## Frequency leakage

If some leakage is acceptable

## Right to be forgotten

Define key name as customer id and utilize json pointers